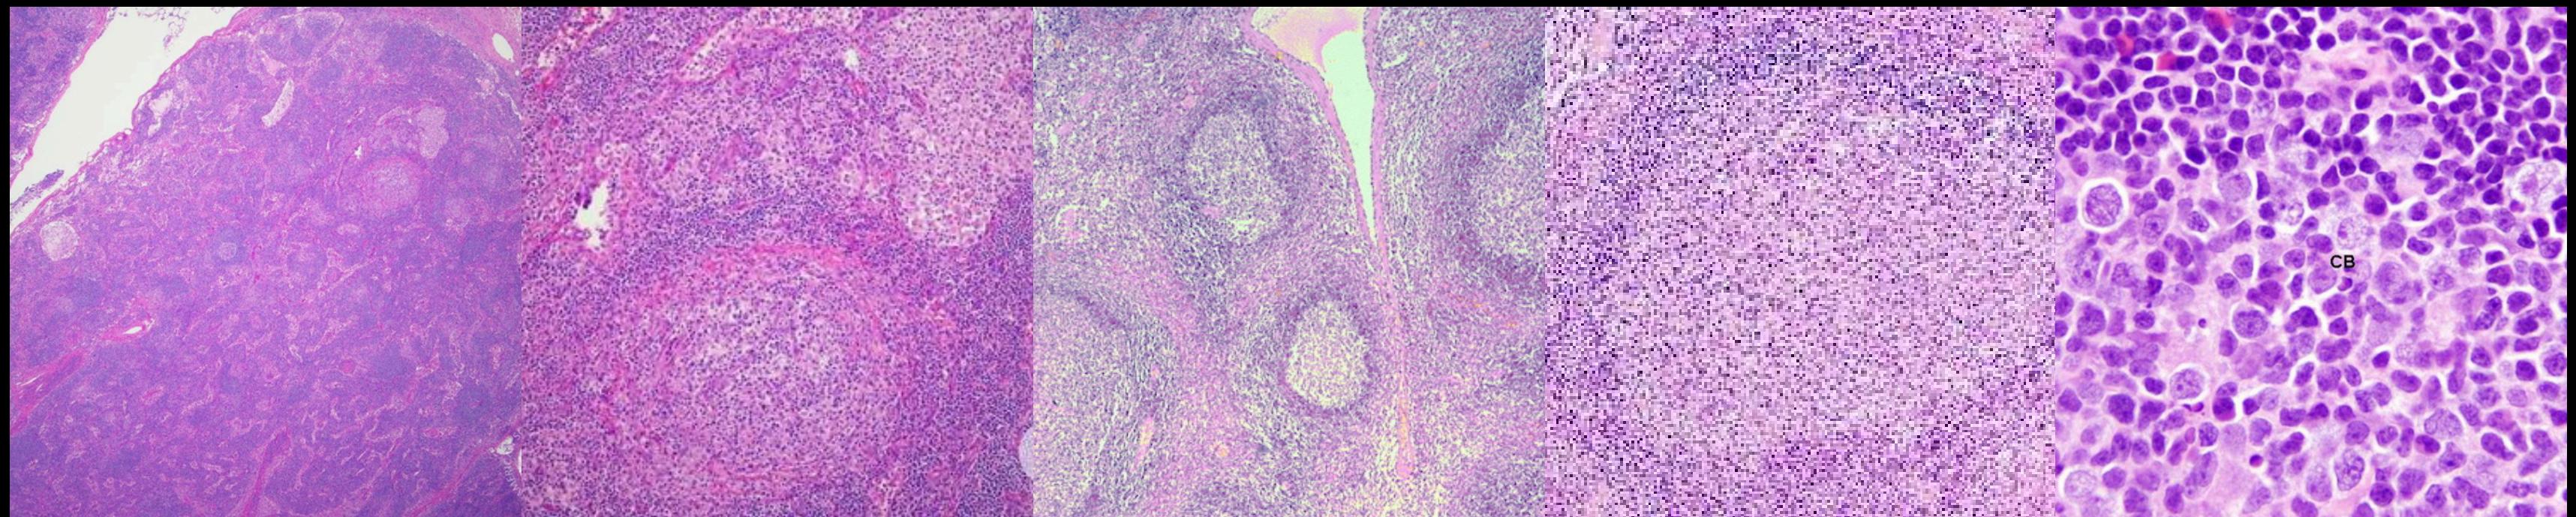


Capstone Project

Metastatic Tissue Detection (Cancer Detection)



Student: Ildar Mamin

Course: INFO-6147-(01)-25F Deep Learning with PyTorch

Date: December 09, 2025

The Problem

- Current Challenge: Pathologists manually examine thousands of tissue slides.
- The Issue: Process is time-consuming, prone to fatigue, and subjective.
- The Need: Automated tools to triage and flag suspicious regions.



Project Objectives

- Primary Goal: Build a Convolutional Neural Network (CNN) to classify lymph node tissue.
- Scope: Binary Classification (Healthy vs. Tumor).
- Deliverable: A functional Web Application for real-time inference.

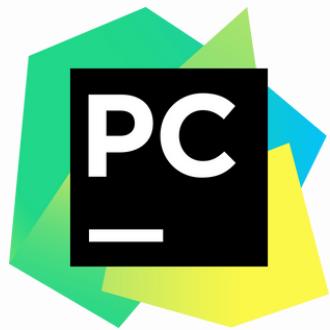
Dataset Description

- Source: PatchCamelyon (PCAM) Benchmark.
- Data Type: H&E stained histopathology images (96x96 pixels).
- Sample Size: Subset of 5,000 images (Balanced distribution).

Source: <https://github.com/basveeling/pcam>

Project structure

Main tools:



Brief Project Structure

- training.ipynb
- app.py
- data

Data Preprocessing

- Resizing: Standardized to 96x96.
- Normalization: Scaled pixel values (Mean=0.5, Std=0.5).
- Splitting: 80% Training / 20% Validation.
- Tensors: Converted to PyTorch format.

```
1 # Data Transformation
2 data_transform = transforms.Compose([
3     transforms.Resize((96, 96)),
4     transforms.ToTensor(),
5     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
6 ])
7
8 print("PCAM Dataset Downloading")
9
10 # Dataset Downloading
11 full_dataset = torchvision.datasets.PCAM(root='./data', split='train', download=True, transform=
12
13 # We need to use subset as a dataset quite big, so I selected the first 5000 imgs.
14 total_images = 5000
15 subset_indices = list(range(total_images))
16 small_dataset = Subset(full_dataset, subset_indices)
17
18 # Splitting Data (80% - training, 20% - val)
19 train_len = int(0.8 * len(small_dataset))
20 val_len = len(small_dataset) - train_len
21
22 train_data, val_data = random_split(small_dataset, [train_len, val_len])
23
24 # Data Loaders - Creation
25 batch_size = 64
26 train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True)
27 val_loader = DataLoader(val_data, batch_size=batch_size, shuffle=False)
28
29 print("Data loaded successfully.")
30 print(f"Training Images: {len(train_data)}")
31 print(f"Validation Images: {len(val_data)}")
32
33 print("Part 2 - Completed")
[11]
```

```
PCAM Dataset Downloading
Downloading...
From (original): https://drive.google.com/uc?id=1Ka0XfEMiwgCYPdTl-vv6eUElOBnKFKQ2
From (redirected): https://drive.usercontent.google.com/download?id=1Ka0XfEMiwgCYPdTl-vv6eUElOBnKFKQ2
To: C:\Users\ilari\PycharmProjects\JupyterNotebook\CapsoneProject\data\pcam\camelyonpatch_level_1_train_jpg_v2_size320x320_label_v2.csv
100%|██████████| 6.42G/6.42G [32:09<00:00, 3.33MB/s]
```

Model Architecture

- Architecture: Custom 3-Layer CNN.
- Feature Extraction: 3 Convolutional Blocks + ReLU + MaxPool.
- Regularization: Dropout (0.5) and Batch Normalization.
- Output: Sigmoid Activation (Probability Score).

```
class MyCNN(nn.Module):  
    def __init__(self):  
        super(MyCNN, self).__init__()  
  
        # First layer: 3 inputs (RGB), 32 outputs  
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)  
        self.bn1 = nn.BatchNorm2d(32)  
        self.pool = nn.MaxPool2d(2, 2)  
  
        # Second layer: 32 inputs, 64 outputs  
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)  
        self.bn2 = nn.BatchNorm2d(64)  
  
        # Third layer: 64 inputs, 128 outputs  
        self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=1)  
        self.bn3 = nn.BatchNorm2d(128)  
  
        # Fully connected layers  
        self.flatten = nn.Flatten()  
        self.fc1 = nn.Linear(128 * 12 * 12, 512)  
        self.dropout = nn.Dropout(0.5) # Dropout to prevent overfitting  
        self.fc2 = nn.Linear(512, 1) # Output is 1 number (probability)  
        self.sigmoid = nn.Sigmoid()  
  
    def forward(self, x):  
        x = self.pool(torch.relu(self.bn1(self.conv1(x))))  
        x = self.pool(torch.relu(self.bn2(self.conv2(x))))  
        x = self.pool(torch.relu(self.bn3(self.conv3(x))))  
  
        x = self.flatten(x)  
        x = torch.relu(self.fc1(x))  
        x = self.dropout(x)  
        x = self.fc2(x)  
        x = self.sigmoid(x)  
        return x  
  
model = MyCNN().to(device)  
  
print("Model Initialized")  
print("Part 3 - Completed")  
[13]  
Model Initialized  
Part 3 - Completed
```

Training & Implementation

- Framework: PyTorch & Python.
- Optimizer: Adam (Learning Rate: 0.001).
- Loss Function: Binary Cross Entropy (BCELoss).
- Epochs: 20

```
# Settings
learning_rate = 0.001
epochs = 20

# Loss and Optimizer
# Binary Cross Entropy Loss (BCELoss)
criterion = nn.BCELoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

# Lists to save values for the graphs
train_losses = []
val_accuracies = []

print("Starting training loop...")
for epoch in range(epochs):
    model.train()
    total_loss = 0.0

    for i, (images, labels) in enumerate(train_loader): # mb
        images = images.to(device)
        # Need to reshape labels to match the output shape
        labels = labels.float().unsqueeze(1).to(device)

        # Forward pass
        outputs = model(images)
        loss = criterion(outputs, labels)

        # Backward pass
        optimizer.zero_grad()
        loss.backward()
        optimizer.step() # optimization

        total_loss += loss.item()

    # Validation step (check accuracy)
    model.eval()
    correct_preds = 0
    total_samples = 0

    with torch.no_grad():
        for images, labels in val_loader:
            images = images.to(device)
            labels = labels.to(device)

            outputs = model(images)
            # If output > 0.5, predict 1 (Tumor), else 0 (Healthy)
            predictions = (outputs > 0.5).float().squeeze()

            total_samples += labels.size(0)
            correct_preds += (predictions == labels).sum().item()

    epoch_loss = total_loss / len(train_loader)
    epoch_acc = 100 * correct_preds / total_samples

    train_losses.append(epoch_loss)
    val_accuracies.append(epoch_acc)

    print(f"Epoch [{epoch+1}/{epochs}], Loss: {epoch_loss:.4f}, Val Accuracy: {epoch_acc:.2f}%")

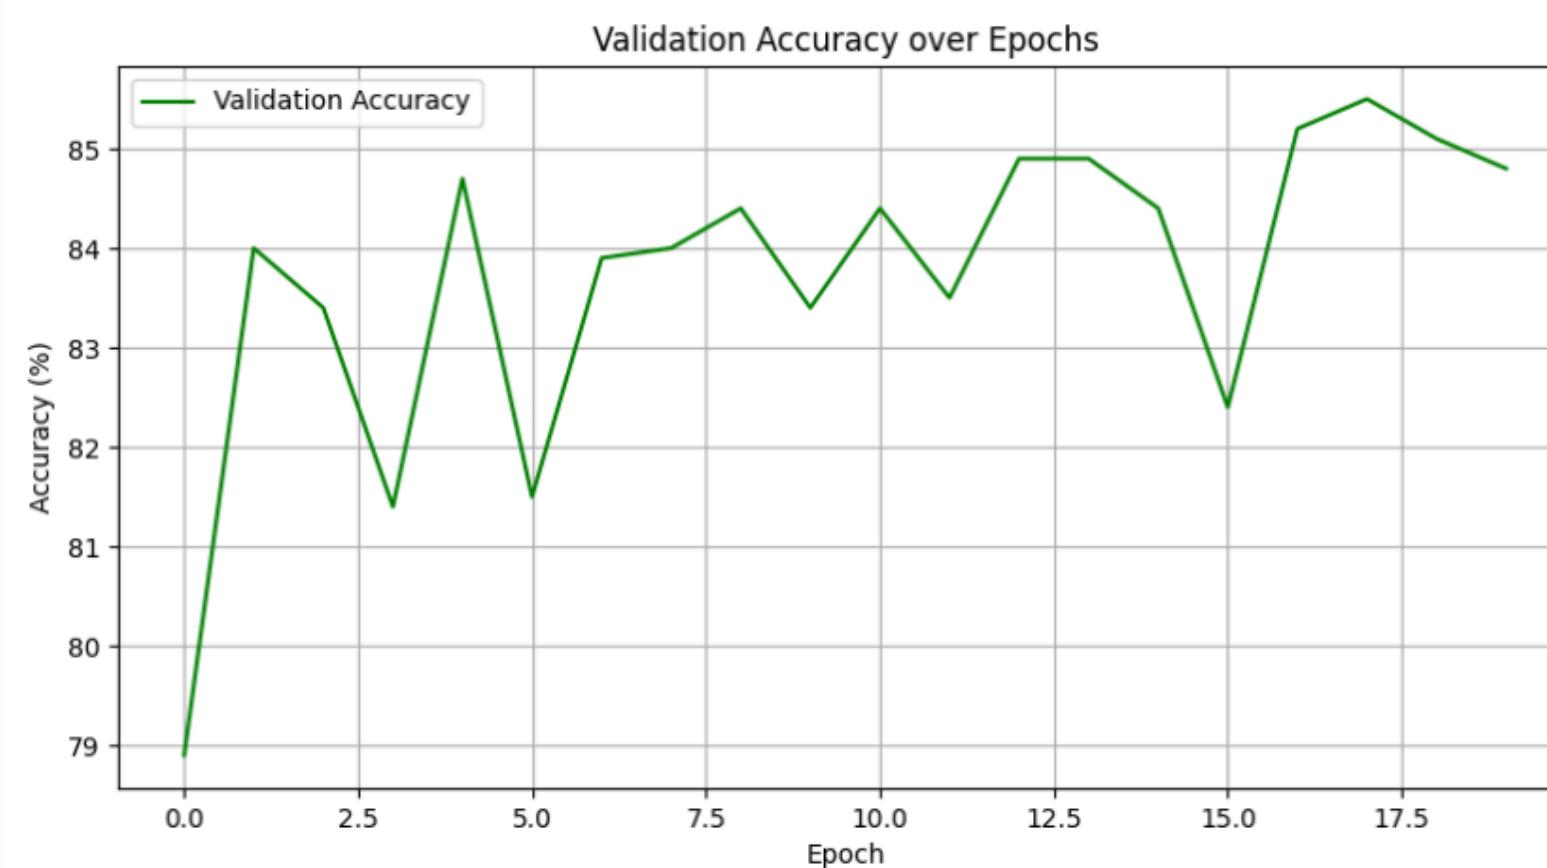
print("Training Finished!")
print("Part 4 - Completed")
[15]
```

Epoch [7/20], Loss: 0.3668, Val Accuracy: 83.90%
Epoch [8/20], Loss: 0.3396, Val Accuracy: 84.00%
Epoch [9/20], Loss: 0.3368, Val Accuracy: 84.40%
Epoch [10/20], Loss: 0.3281, Val Accuracy: 83.40%
Epoch [11/20], Loss: 0.3172, Val Accuracy: 84.40%
Epoch [12/20], Loss: 0.3181, Val Accuracy: 83.50%
Epoch [13/20], Loss: 0.3059, Val Accuracy: 84.90%
Epoch [14/20], Loss: 0.3069, Val Accuracy: 84.90%
Epoch [15/20], Loss: 0.2781, Val Accuracy: 84.40%
Epoch [16/20], Loss: 0.2681, Val Accuracy: 82.40%
Epoch [17/20], Loss: 0.2617, Val Accuracy: 85.20%
Epoch [18/20], Loss: 0.2788, Val Accuracy: 85.50%
Epoch [19/20], Loss: 0.2523, Val Accuracy: 85.10%
Epoch [20/20], Loss: 0.2323, Val Accuracy: 84.80%

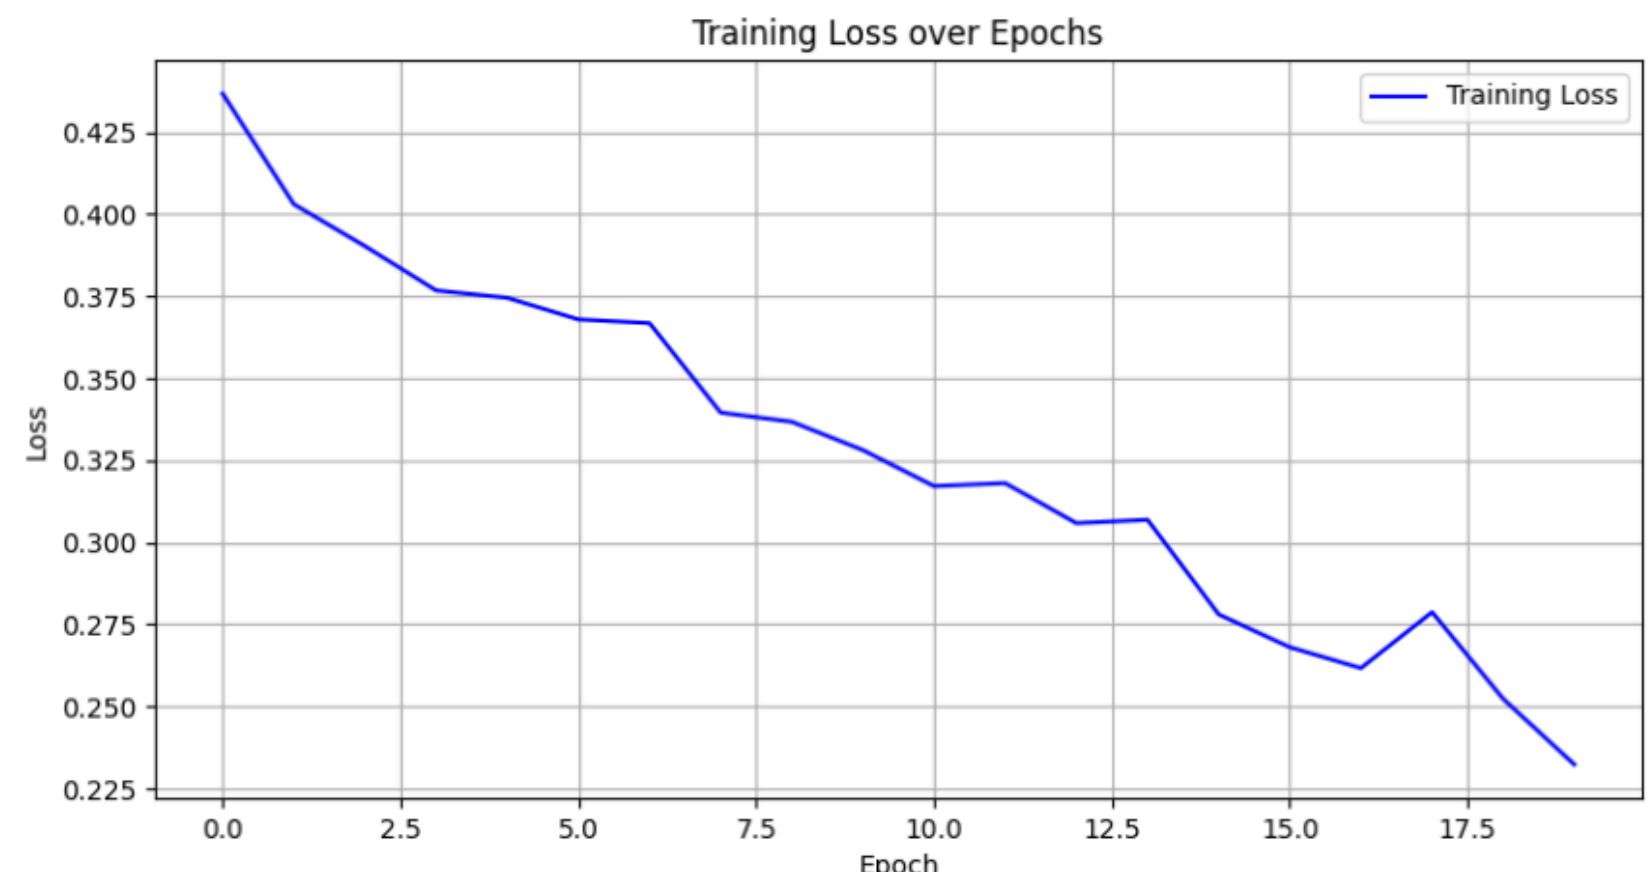
Training Finished!
Part 4 - Completed

Results & Evaluation

Accuracy

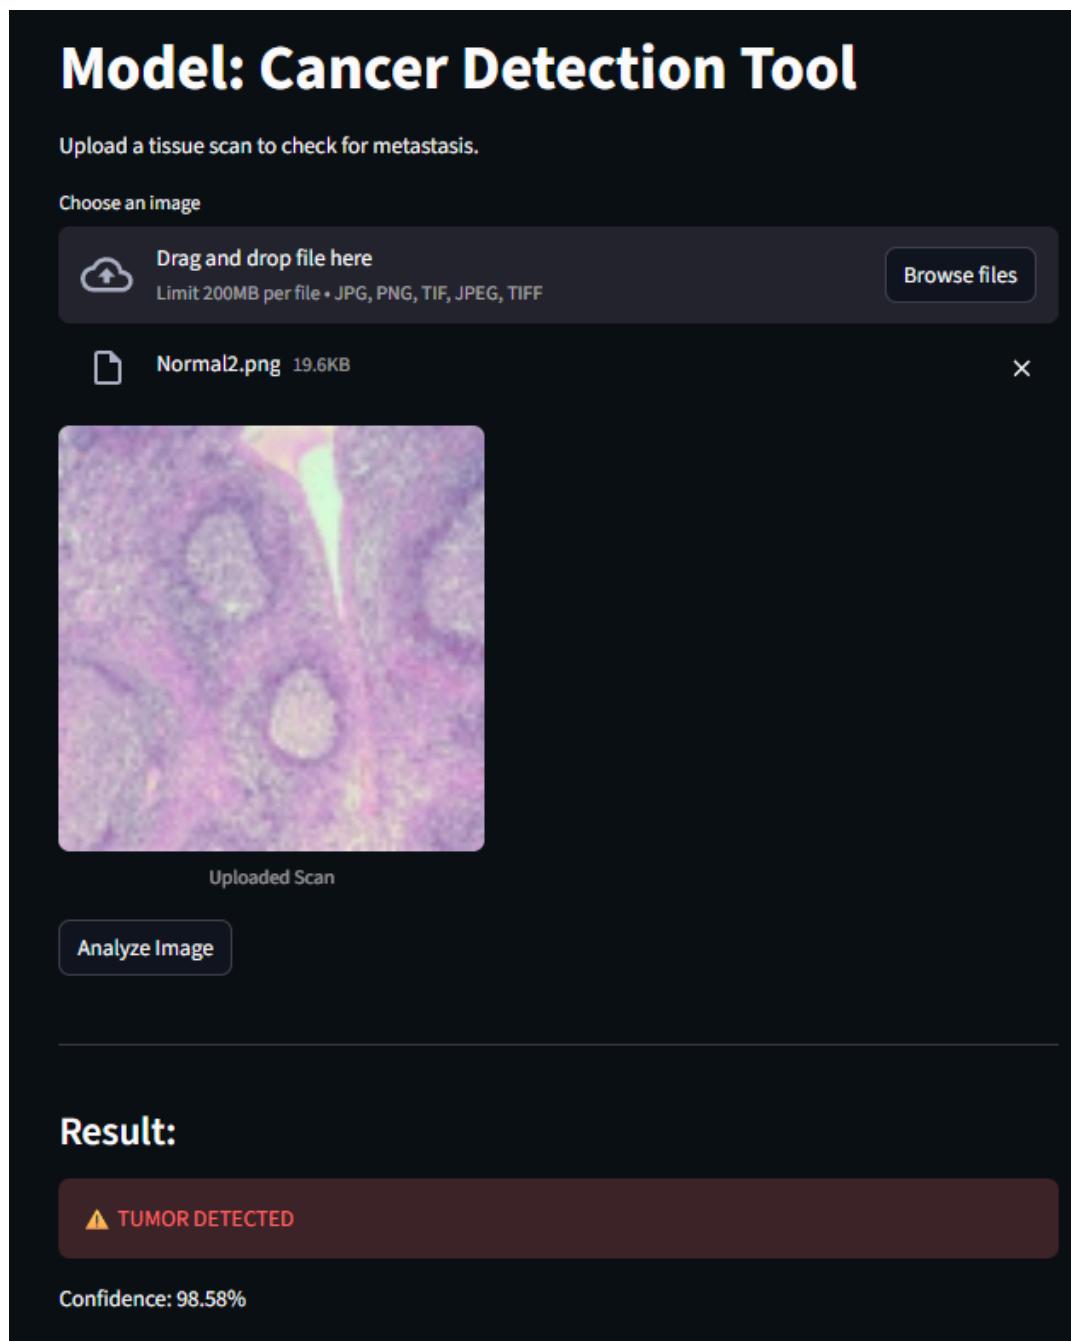


Loss

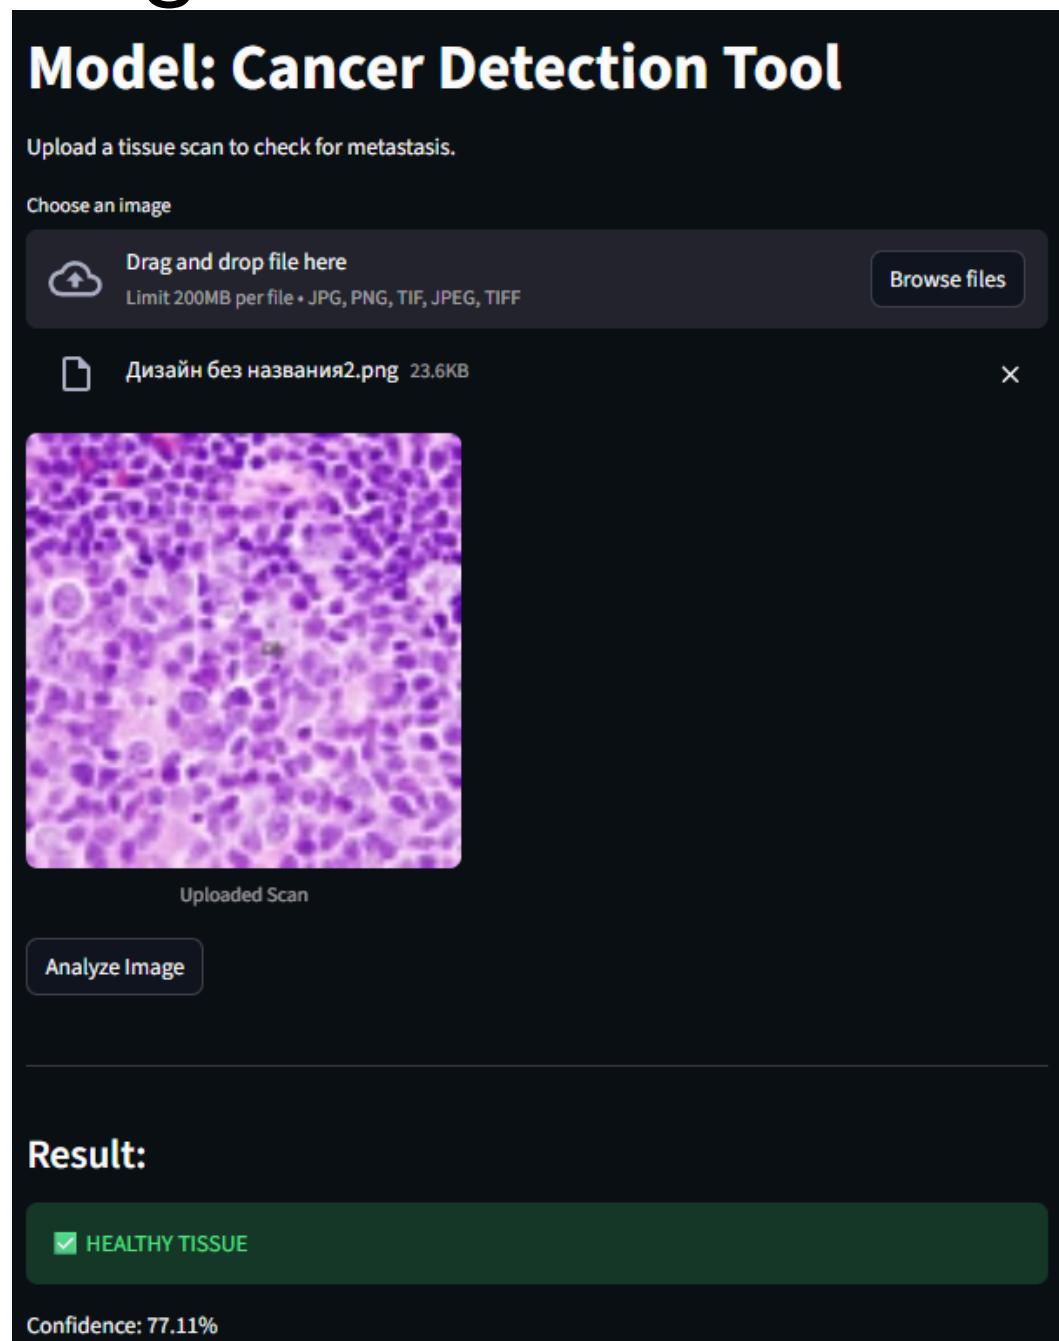


Deployment (The App)

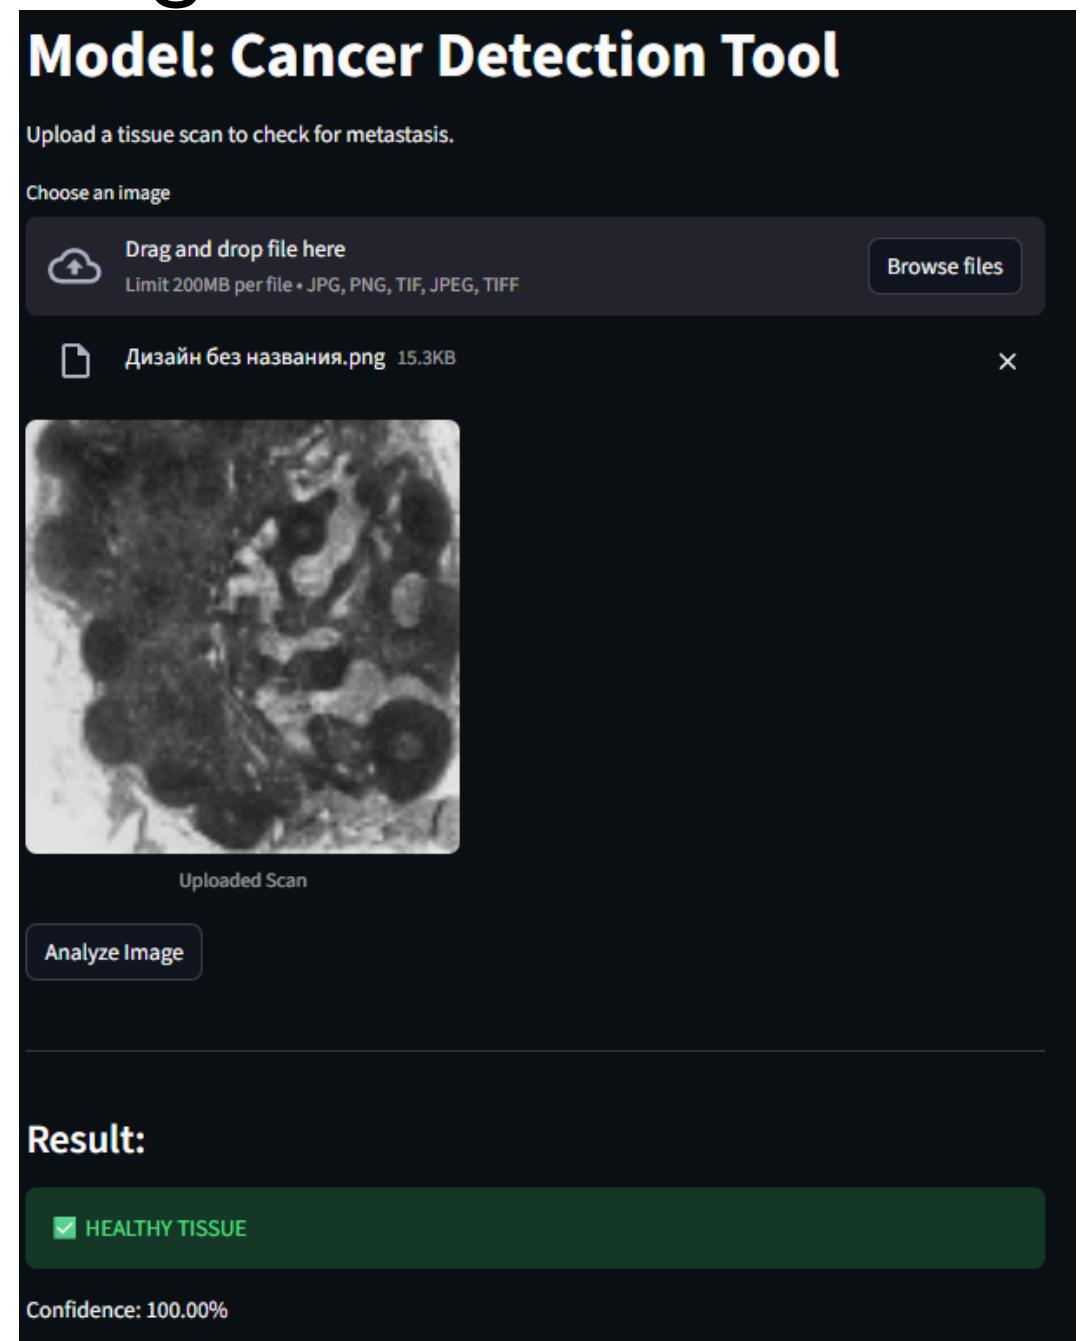
Possitive acc 98%



Neg acc 77%



Neg acc 99%



Conclusion

- Conclusion: Successfully automated metastasis detection with high accuracy.
- Limitations: Trained on a subset; no heavy data augmentation.

Possible Improvements:

- Train on full dataset.
- Implement Transfer Learning (ResNet/VGG).
- Advanced Augmentation (Rotation/Flipping)